


On Quantum Chosen-Ciphertext Attacks and Learning with Errors

Gorjan Alagic[1], Stacey Jeffery[2], Maris Ozols[3], and Alexander Poremba[4]

[1] QuICS, University of Maryland, and NIST, Gaithersburg, MD, USA
[2] QuSoft and CWI, Amsterdam, Netherlands
[3] QuSoft and University of Amsterdam, Amsterdam, Netherlands
[4] Computing and Mathematical Sciences, Caltech, Pasadena, CA, USA

arXiv:1808.09655v2 [quant-ph] 23 Jun 2019

Abstract. Large-scale quantum computing is a significant threat to classical public-key cryptography. In strong "quantum access" security models, numerous symmetric-key cryptosystems are also vulnerable. We consider classical encryption in a model which grants the adversary quantum oracle access to encryption and decryption, but where the latter is restricted to non-adaptive (i.e., pre-challenge) queries only. We define this model formally using appropriate notions of ciphertext indistinguishability and semantic security (which are equivalent by standard arguments) and call it QCCA1 in analogy to the classical CCA1 security model. Using a bound on quantum random-access codes, we show that the standard PRF- and PRP-based encryption schemes are QCCA1-secure when instantiated with quantum-secure primitives.

We then revisit standard IND-CPA-secure Learning with Errors (LWE) encryption and show that leaking just one quantum decryption query (and no other queries or leakage of any kind) allows the adversary to recover the full secret key with constant success probability. In the classical setting, by contrast, recovering the key uses a linear number of decryption queries, and this is optimal. The algorithm at the core of our attack is a (large-modulus version of) the well-known Bernstein-Vazirani algorithm. We emphasize that our results should **not** be interpreted as a weakness of these cryptosystems in their stated security setting (i.e., post-quantum chosen-plaintext secrecy). Rather, our results mean that, if these cryptosystems are exposed to chosen-ciphertext attacks (e.g., as a result of deployment in an inappropriate real-world setting) then quantum attacks are even more devastating than classical ones.

1 Introduction

1.1 Background

Large-scale quantum computers pose a dramatic threat to classical cryptography. The ability of such devices to run Shor's efficient quantum factoring algorithm (and its variants) would lead to devastation of the currently deployed public-key cryptography infrastructure [CJL+16; Sho94]. This threat has led to significant work on so-called "post-quantum" alternatives, where a prominent category is occupied by cryptosystems based on the *Learning with Errors* (LWE) problem of solving noisy linear equations over $\mathbb{Z}_q$ [Reg05] and its variants [CJL+16; NIS17].

In addition to motivating significant work on post-quantum cryptosystems, the threat of quantum computers has also spurred general research on secure classical cryptography in the presence of quantum adversaries. One area in particular explores strong security models where a quantum adversary gains precise quantum control over portions of a classical cryptosystem. In such models, a number of basic symmetric-key primitives can be broken by simple quantum attacks based on Simon's algorithm [KM10; KM12; KLL+16; SS17; Sim97]. It is unclear if the assumption behind these models is plausible for typical physical implementations of symmetric-key cryptography. However, attacks which involve quantumly querying a classical function are always available in scenarios where the adversary has access to a circuit for the relevant function. This is the case for hashing, public-key encryption, and circuit obfuscation. Moreover, understanding this model is crucial for gauging the degree to which any physical device involved in cryptography must be resistant to