

2015–
2019

# Quantum-inspired low-rank stochastic regression with logarithmic dependence on the dimension

András Gilyén[1,*] Seth Lloyd[2] Ewin Tang[3]

November 13, 2018

arXiv:1811.04909v1 [cs.DS] 12 Nov 2018

**Abstract**

We construct an efficient classical analogue of the quantum matrix inversion algorithm [HHL09] for low-rank matrices. Inspired by recent work of Tang [Tan18a], assuming length-square sampling access to input data, we implement the pseudoinverse of a low-rank matrix and sample from the solution to the problem $Ax = b$ using fast sampling techniques. We implement the pseudo-inverse by finding an approximate singular value decomposition of $A$ via subsampling, then inverting the singular values. In principle, the approach can also be used to apply any desired "smooth" function to the singular values. Since many quantum algorithms can be expressed as a singular value transformation problem [GSLW18], our result suggests that more low-rank quantum algorithms can be effectively "dequantised" into classical length-square sampling algorithms.

## 1 Introduction

Quantum computing provides potential exponential speed-ups over classical computers for a variety of linear algebraic tasks, including an operational version of matrix inversion [HHL09]. Recently, inspired by the quantum algorithm for recommendation systems [KP17], Tang showed how to generalize the well-known FKV algorithm [FKV04] to sample from the singular vectors of low-rank matrices [Tan18a] and to implement principal component component analysis [Tan18b]. Intriguingly, Tang's work suggests that many of the quantum algorithms for low-rank matrix manipulation [RSML18] can be extended to provide fast *classical* algorithms under suitable sampling assumptions, achieving logarithmic dependence on the dimension. In this work, we show that such exponential speed-ups are indeed possible in the case of low-rank matrix inversion. Our treatment is self-contained and improves some aspects of previous approaches [Tan18a], leading to smaller exponents in our runtime bounds.

Suppose we want to solve $Ax = b$, where we are given $A \in \mathbb{R}^{m\times n}$ and $b \in \mathbb{R}^m$, and wish to recover $x \in \mathbb{R}^n$. The equation might not have a solution, but we can always find an $x$ minimizing $\|Ax - b\|$. Namely, $x = A^+b$ works, where $A^+$ is the pseudoinverse of $A$. If

$$A = \sum_{\ell=1}^{k} \sigma_\ell \, u^{(\ell)} v^{(\ell)^T}$$

is a singular value decomposition of $A$, such that $\sigma_\ell > 0$, then the pseudoinverse is simply $A^+ = \sum_{\ell=1}^{k} v^{(\ell)} u^{(\ell)^T}/\sigma_\ell$, and $x = \sum_{\ell=1}^{k} v^{(\ell)} \langle u^{(\ell)}, b\rangle/\sigma_\ell$. The problem of finding the pseudoinverse of

[1]QuSoft, CWI and University of Amsterdam, the Netherlands. Supported by ERC Consolidator Grant QPROGRESS and partially supported by QuantERA project QuantAlgo 680-91-034.

[2]Massachusetts Institute of Technology, Departments of Mechanical Engineering and Physics, Xanadu
[3]University of Washington
[*]Corresponding author: gilyen@cwi.nl