


2017–
2021



Security of the Fiat-Shamir Transformation in the Quantum Random-Oracle Model*

Jelle Don[1,2], Serge Fehr[1,3,4], Christian Majenz[2,4], and Christian Schaffner[2,4]

[1] Centrum Wiskunde & Informatica (CWI), Amsterdam, Netherlands
[2] Institute for Logic, Language and Computation, University of Amsterdam, Amsterdam, Netherlands
[3] Mathematical Institute, Leiden University, Netherlands
[4] QuSoft, Amsterdam, Netherlands
jelle.don@cwi.nl, serge.fehr@cwi.nl, c.majenz@uva.nl, c.schaffner@uva.nl

arXiv:1902.07556v4 [cs.CR] 27 Jul 2020

Abstract. The famous Fiat-Shamir transformation turns any public-coin three-round interactive proof, i.e., any so-called Σ-protocol, into a non-interactive proof in the random-oracle model. We study this transformation in the setting of a *quantum adversary* that in particular may query the random oracle in quantum superposition.

Our main result is a generic reduction that transforms any quantum dishonest prover attacking the Fiat-Shamir transformation in the quantum random-oracle model into a similarly successful quantum dishonest prover attacking the underlying Σ-protocol (in the standard model). Applied to the standard soundness and proof-of-knowledge definitions, our reduction implies that both these security properties, in both the computational and the statistical variant, are preserved under the Fiat-Shamir transformation even when allowing quantum attacks. Our result improves and completes the partial results that have been known so far, but it also proves wrong certain claims made in the literature.

In the context of post-quantum secure signature schemes, our results imply that for any Σ-protocol that is a proof-of-knowledge against quantum dishonest provers (and that satisfies some additional natural properties), the corresponding Fiat-Shamir signature scheme is secure in the quantum random-oracle model. For example, we can conclude that the non-optimized version of Fish, which is the bare Fiat-Shamir variant of the NIST candidate Picnic, is secure in the quantum random-oracle model.

1 Introduction

The (quantum) random-oracle model. The *random-oracle model* (ROM) is a means to treat a cryptographic hash function H as an ideal primitive. In the ROM, the only way to "compute" the hash $H(x)$ of any value x is by making a *query* to an imaginary entity, the *random oracle* (RO), which has chosen H uniformly at random from the set of *all* functions with the considered domain and range.

The hope is that if a cryptographic scheme is secure in the ROM then it is also secure in the standard model, as long as H is instantiated with a "good enough" cryptographic hash function. Even though in general we cannot hope to obtain provable security in the standard model in this way, (since there exist artificial counter examples [CGH04]), this approach works extremely well in practice, leading to very efficient schemes that tend to resist all known attacks.

What makes the ROM particularly convenient is that in the security proof of a cryptographic scheme, we can control the RO. For instance, simply by recording the queries that the adversary makes to the RO, we know exactly which hash values he knows, and the hash value $H(x)$ is random to him for any x that he has not queried. Furthermore, we can *reprogram* the RO, meaning that we can let $H(x)$ be some particular value y for some specific x, as long as it is random from the adversary's perspective.

When considering a *quantum* adversary, the picture changes a bit. In order to model that such an adversary can evaluate any hash function *in superposition* on different inputs, we must allow such a quantum adversary in the ROM to make superposition queries to the RO: for any superposition $\sum_x \alpha_x|x\rangle$ it may

* ©IACR 2019. This is a reformatted and slightly modified version of the article submitted by the authors to the IACR and to Springer-Verlag in May 2019. The published version is available from the proceedings of CRYPTO 2019.