
2016–
2022


2015–
2019



Quantum algorithms for zero-sum games

Joran van Apeldoorn* András Gilyén†

QuSoft, CWI and University of Amsterdam, the Netherlands.
e-mail: {apeldoor,gilyen}@cwi.nl

April 8, 2019

arXiv:1904.03180v1 [quant-ph] 5 Apr 2019

Abstract

We derive sublinear-time quantum algorithms for computing the Nash equilibrium of two-player zero-sum games, based on efficient Gibbs sampling methods. We are able to achieve speed-ups for both dense and sparse payoff matrices at the cost of a mildly increased dependence on the additive error compared to classical algorithms. In particular we can find ε-approximate Nash equilibrium strategies in complexity $\widetilde{\mathcal{O}}\big(\sqrt{n+m}/\varepsilon^3\big)$ and $\widetilde{\mathcal{O}}\big(\sqrt{s}/\varepsilon^{3.5}\big)$ respectively, where $n \times m$ is the size of the matrix describing the game and s is its sparsity. Our algorithms use the LP formulation of the problem and apply techniques developed in recent works on quantum SDP-solvers. We also show how to reduce general LP-solving to zero-sum games, resulting in quantum LP-solvers that have complexities $\widetilde{\mathcal{O}}\big(\sqrt{n+m}\gamma^3\big)$ and $\widetilde{\mathcal{O}}\big(\sqrt{s}\gamma^{3.5}\big)$ for the dense and sparse access models respectively, where γ is the relevant "scale-invariant" precision parameter.

1 Introduction

A *matrix game* is a two-player zero-sum game in which both players have only a finite number of pure strategies. We label the moves for the first player (called Alice) by $[n]$ and the moves for the second player (called Bob) by $[m]$. If Alice plays $i \in [n]$ and Bob plays $j \in [m]$, then Alice gets a *payoff* $A_{ij} \in [-1, 1]$ and Bob gets payoff $-A_{ij}$. Their individual goal is to get the highest payoff possible. The payoff can be written in matrix form $A \in [-1, 1]^{n\times m}$, hence the name *matrix game*. A game is called *symmetric* if $m = n$ and $A = -A^T$, in other words, the payoff matrix is *skew symmetric*.

If one of the players would always play the same move, then (for most games) it would be easy for the other player to win. Hence a strategy will be randomized in general. Let Δ^n be the set of all non-negative vectors in $\mathbb{R}^n$ that sum to 1, i.e., the set of all probability distributions over n elements. Given a randomized strategy $x \in \Delta^n$ for Alice and a randomized strategy $y \in \Delta^m$ for Bob, the expected payoff for Alice is $x^T Ay$.

Naturally Bob's goal is to minimize Alice's expected payoff; the best he can do is to assume that Alice plays the best strategy x on her side and optimize his y for that:

$$\min_{y\in\Delta^m} \max_{x\in\Delta^n} x^T Ay$$

*Supported by the Netherlands Organization for Scientific Research, grant number 617.001.351 and partially supported by QuantERA project QuantAlgo 680-91-034.

†Supported by ERC Consolidator Grant 615307-QPROGRESS and partially supported by QuantERA project QuantAlgo 680-91-034.