

2019–
2021

Towards Quantum Advantage via Topological Data Analysis

Casper Gyurik[1], Chris Cade[2], and Vedran Dunjko[1,3]

[1]LIACS, Leiden University, Niels Bohrweg 1, 2333 CA Leiden, Netherlands
[2]QuSoft, Centrum Wiskunde & Informatica (CWI), Science Park 123, 1098 XG Amsterdam, Netherlands
[3]LION, Leiden University, Niels Bohrweg 2, 2333 CA Leiden, Netherlands
August 28th, 2022

arXiv:2005.02607v5 [quant-ph] 21 Oct 2022

Even after decades of quantum computing development, examples of generally useful quantum algorithms with exponential speedups over classical counterparts are scarce. Recent progress in quantum algorithms for linear-algebra positioned quantum machine learning (QML) as a potential source of such useful exponential improvements. Yet, in an unexpected development, a recent series of "dequantization" results has equally rapidly removed the promise of exponential speedups for several QML algorithms. This raises the critical question whether exponential speedups of other linear-algebraic QML algorithms persist. In this paper, we study the quantum-algorithmic methods behind the algorithm for topological data analysis of Lloyd, Garnerone and Zanardi through this lens. We provide evidence that the problem solved by this algorithm is classically intractable by showing that its natural generalization is as hard as simulating the one clean qubit model – which is widely believed to require superpolynomial time on a classical computer – and is thus very likely immune to dequantizations. Based on this result, we provide a number of new quantum algorithms for problems such as rank estimation and complex network analysis, along with complexity-theoretic evidence for their classical intractability. Furthermore, we analyze the suitability of the proposed quantum algorithms for near-term implementations. Our results provide a number of useful applications for full-blown, and restricted quantum computers with a guaranteed exponential speedup over classical methods, recovering some of the potential for linear-algebraic QML to become one of quantum computing's killer applications.

1 Introduction

Quantum machine learning (QML) is a rapidly growing field [1, 2] that has brought forth numerous proposals regarding ways for quantum computers to help analyze data. Several of these proposals involve using quantum algorithms for linear algebra – most notably Harrow, Hassidim and Lloyd's matrix inversion algorithm [3] – to exponentially speed up tasks in machine learning. Other proposals such as the use of parameterized quantum circuits [4, 5, 6] provide a different approach based on identifying genuinely new quantum learning models (rather than speedups of established methods), which are more amenable to near-term quantum computing restrictions. These QML proposals have all been hailed as possible examples of quantum computing's "killer application": genuinely and broadly useful quantum algorithms which superpolynomially outperform their best known classical counterparts (which are very rare even if full-blown quantum computing is assumed).

However, previously speculated superpolynomial speedups of linear-algebraic QML proposals were revealed to actually be at most polynomial speedups, as exponentially faster classical algorithms were devised that operate under analogous assumptions [7, 8]. Nevertheless, practically relevant polynomial speedups may persist [9, 10]. While quadratic speedups have obvious appeal on paper, recent analysis involving concrete near-term device properties revealed that low-degree

Casper Gyurik: c.f.s.gyurik@liacs.leidenuniv.nl