

arXiv:2009.08360v1 [quant-ph] 17 Sep 2020

Improved Quantum Boosting

Adam Izdebski* Ronald de Wolf†

September 18, 2020

Abstract

Boosting is a general method to convert a weak learner (which generates hypotheses that are just slightly better than random) into a strong learner (which generates hypotheses that are much better than random). Recently, Arunachalam and Maity [AM20] gave the first quantum improvement for boosting, by combining Freund and Schapire's AdaBoost algorithm with a quantum algorithm for approximate counting. Their booster is faster than classical boosting as a function of the VC-dimension of the weak learner's hypothesis class, but worse as a function of the quality of the weak learner. In this paper we give a substantially faster and simpler quantum boosting algorithm, based on Servedio's SmoothBoost algorithm [Ser03].

1 Introduction

1.1 Boosting

There has been tremendous growth in machine learning research and applications, both in practice (applying all sorts of methods on all sorts of data and seeing what works well) and in theory (computational learning theory). However, not very many ideas generated in theoretical machine learning have had a large impact on machine learning practice. One of the exceptions is *boosting*, which is a simple, general, and widely applicable method to improve the generalization error of a given learning method, i.e., to convert a *weak* learner into a *strong* learner.

The set-up here is binary classification: we are trying to predict binary labels y from points $x \in \mathcal{X}$. A typical case would be $\mathcal{X} = \{0,1\}^n$. We are given m labeled examples $(x_1, y_1), \ldots, (x_m, y_m) \in \mathcal{X} \times \{-1, 1\}$ where the x_is are independent and identically distributed (i.i.d.) according to some unknown distribution $\mathcal{D}$, and the binary labels are determined by some unknown *target function* $f : \mathcal{X} \to \{-1, 1\}$ that we are trying to learn, i.e., $y_i = f(x_i)$. A weak learner $\mathcal{W}$ is an algorithm that can be fed a number of examples according to a specified distribution D over the m examples of the given sample, and that is then promised to generate a hypothesis $h : \mathcal{X} \to \{-1, 1\}$ that is slightly better than random w.r.t. that D:

$$\Pr_{x \sim D}[h(x) \neq f(x)] \leq 1/2 - \gamma.$$

*Work done while a student at the University of Amsterdam. adam.izdebski1@gmail.com

†QuSoft, CWI and University of Amsterdam, the Netherlands. Partially supported by ERC Consolidator Grant 615307-QPROGRESS (which ended February 2019), and by the Dutch Research Council (NWO) through Gravitation-grant Quantum Software Consortium, 024.003.037, and through QuantERA project QuantAlgo 680-91-034. rdewolf@cwi.nl

