

Lightweight Detection of a Small Number of Large Errors in a Quantum Circuit

Noah Linden[1] and Ronald de Wolf[2]

[1]School of Mathematics, University of Bristol. n.linden@bristol.ac.uk

[2]QuSoft, CWI and University of Amsterdam, the Netherlands. rdewolf@cwi.nl

arXiv:2009.08840v3 [quant-ph] 13 Apr 2021

Suppose we want to implement a unitary U, for instance a circuit for some quantum algorithm. Suppose our actual implementation is a unitary $\tilde{U}$, which we can only apply as a black-box. In general it is an exponentially-hard task to decide whether $\tilde{U}$ equals the intended U, or is significantly different in a worst-case norm. In this paper we consider two special cases where relatively efficient and lightweight procedures exist for this task.

First, we give an efficient procedure under the assumption that U and $\tilde{U}$ (both of which we can now apply as a black-box) are either equal, or differ significantly in only *one* k-qubit gate, where $k = O(1)$ (the k qubits need not be contiguous). Second, we give an even more lightweight procedure under the assumption that U and $\tilde{U}$ are *Clifford* circuits which are either equal, or different in arbitrary ways (the specification of U is now classically given while $\tilde{U}$ can still only be applied as a black-box). Both procedures only need to run $\tilde{U}$ a constant number of times to detect a constant error in a worst-case norm. We note that the Clifford result also follows from earlier work of Flammia and Liu [FL11] and da Silva, Landon-Cardinal, and Poulin [dSLCP11].

In the Clifford case, our error-detection procedure also allows us efficiently to *learn* (and hence correct) $\tilde{U}$ if we have a small list of possible errors that could have happened to U; for example if we know that only $O(1)$ of the gates of $\tilde{U}$ are wrong, this list will be polynomially small and we can test each possible erroneous version of U for equality with $\tilde{U}$.

1 Introduction

With the first tentative steps for implementing quantum computations on larger numbers of qubits (53 qubits in the case of Google's quantum supremacy experiment [AM19]) comes the need to *verify* whether those implementations actually work as intended. In contrast to classical computations, we cannot just "open up" the computer midway through the computation to check whether everything is still on track and then allow the computation to continue, because measurements on the intermediate quantum state typically destroy the superposition; and learning the quantum state takes exponential effort in the number of qubits in general. Similarly, simulation of a general n-qubit quantum circuit to determine what the intended output should be on a given input state, becomes infeasible if $n \gg 50$.

Even reasonably-well implemented circuits of simple quantum operations ("gates") can still be marred by many different types of errors: a few large errors (where a gate or qubit is totally wrong or even absent), or many smallish errors (for example slight overrotations), or some combination of both. Strategies are needed to deal with these. In the long run,