
2016–
2022


2019–
2023





Quantum algorithms for matrix scaling and matrix balancing

Joran van Apeldoorn* Sander Gribling† Yinan Li‡ Harold Nieuwboer§
Michael Walter¶ Ronald de Wolf‖

arXiv:2011.12823v1 [quant-ph] 25 Nov 2020

Abstract

Matrix scaling and matrix balancing are two basic linear-algebraic problems with a wide variety of applications, such as approximating the permanent, and pre-conditioning linear systems to make them more numerically stable. We study the power and limitations of quantum algorithms for these problems.

We provide quantum implementations of two classical (in both senses of the word) methods: Sinkhorn's algorithm for matrix scaling and Osborne's algorithm for matrix balancing. Using amplitude estimation as our main tool, our quantum implementations both run in time $\widetilde{O}(\sqrt{mn}/\varepsilon^4)$ for scaling or balancing an $n \times n$ matrix (given by an oracle) with m non-zero entries to within ℓ_1-error ε. Their classical analogs use time $\widetilde{O}(m/\varepsilon^2)$, and every classical algorithm for scaling or balancing with small constant ε requires $\Omega(m)$ queries to the entries of the input matrix. We thus achieve a polynomial speed-up in terms of n, at the expense of a worse polynomial dependence on the obtained ℓ_1-error ε. We emphasize that even for constant ε these problems are already non-trivial (and relevant in applications).

Along the way, we extend the classical analysis of Sinkhorn's and Osborne's algorithm to allow for errors in the computation of marginals. We also adapt an improved analysis of Sinkhorn's algorithm for entrywise-positive matrices to the ℓ_1-setting, leading to an $\widetilde{O}(n^{1.5}/\varepsilon^3)$-time quantum algorithm for ε-ℓ_1-scaling in this case.

We also prove a lower bound, showing that our quantum algorithm for matrix scaling is essentially optimal for constant ε: every quantum algorithm for matrix scaling that achieves a constant ℓ_1-error with respect to uniform marginals needs to make at least $\Omega(\sqrt{mn})$ queries.

*Institute for Information Law and QuSoft, University of Amsterdam, the Netherlands. Supported by the Dutch Research Council (NWO/OCW), as part of the Quantum Software Consortium programme (project number 024.003.037). work@bitofbytes.com

†IRIF, Université de Paris, CNRS, Paris, France. Partially supported by SIRTEQ-grant QuIPP. gribling@irif.fr

‡Graduate School of Mathematics, Nagoya University, Japan. Partially supported by MEXT Quantum Leap Flagship Program (MEXT Q-LEAP) Grant Number JPMXS0120319794. Yinan.Li@math.nagoya-u.ac.jp

§Korteweg–de Vries Institute for Mathematics and QuSoft, University of Amsterdam. h.a.nieuwboer@uva.nl

¶Korteweg–de Vries Institute for Mathematics, Institute for Theoretical Physics, Institute for Language, Logic, and Computation, and QuSoft, University of Amsterdam. Partially supported by the Dutch Research Council (NWO) through Veni grant no. 680-47-459. m.walter@uva.nl

‖QuSoft, CWI and University of Amsterdam, the Netherlands. Partially supported by the Dutch Research Council (NWO/OCW), as part of the Quantum Software Consortium programme (project number 024.003.037), and through QuantERA ERA-NET Cofund project QuantAlgo (680-91-034). rdewolf@cwi.nl