


2017–
2021



Online-Extractability in the Quantum Random-Oracle Model

Jelle Don[1], Serge Fehr[1,2], Christian Majenz[1,4], and Christian Schaffner[3,4]

[1] Centrum Wiskunde & Informatica (CWI), Amsterdam, Netherlands
[2] Mathematical Institute, Leiden University, Netherlands
[3] Institute for Logic, Language and Computation, University of Amsterdam, Amsterdam, Netherlands
[4] QuSoft, Amsterdam, Netherlands
jelle.don@cwi.nl, serge.fehr@cwi.nl, christian.majenz@cwi.nl, c.schaffner@uva.nl

arXiv:2103.03085v2 [cs.CR] 17 Sep 2021

Abstract. We show the following generic result. Whenever a quantum query algorithm in the quantum random-oracle model outputs a classical value t that is promised to be in some tight relation with $H(x)$ for some x, then x can be efficiently extracted with almost certainty. The extraction is by means of a suitable simulation of the random oracle and works *online*, meaning that it is *straightline*, i.e., without rewinding, and *on-the-fly*, i.e., during the protocol execution and without disturbing it.

The technical core of our result is a new commutator bound that bounds the operator norm of the commutator of the unitary operator that describes the evolution of the compressed oracle (which is used to simulate the random oracle above) and of the measurement that extracts x.

We show two applications of our generic online extractability result. We show *tight* online extractability of commit-and-open Σ-protocols in the quantum setting, and we offer the first complete post-quantum security proof of the *textbook* Fujisaki-Okamoto transformation, i.e, without adjustments to facilitate the proof, including concrete security bounds.

1 Introduction

Background. *Extractability* plays an important role in cryptography. In an extractable protocol, on a high level, an algorithm $\mathcal{A}$ sends messages that depend on some secret s, and while the secret remains private in an honest run of the protocol, an *extractor* can learn s via some form of enhanced access to $\mathcal{A}$. The probably most prominent example is that of (zero-knowledge) *proofs* (or *arguments*) *of knowledge*, for which, by definition, there must exist an extractor that manages to extract a witness from any successful yet possibly dishonest prover. Another example are *extractable commitments*, which have a wide range of applications. Hash-based extractable commitments are extremely simple to construct and prove secure in the random-oracle model (ROM) [Pas03]. Indeed, when the considered hash function H is modelled as a random oracle, the hash input x for the commitment $c = H(x)$, where $x = s\|r$ consists of the actual secret s and randomness r, can be extracted simply by finding a query x to the random oracle that yielded c as an output.

The general notion of extractability comes in different flavors. The most well-known example is extraction by *rewinding*. Here, the extractor is allowed to run $\mathcal{A}$ several times, on the same private input and using different randomness. This is the notion usually considered in the context of proofs/arguments of knowledge. In some contexts, extraction via rewinding access is not possible. For example, the UC security model prohibits the simulator to rewind the adversary. In other occasions, rewinding may be possible but not desirable due to a loss of efficiency, which stems from having to run $\mathcal{A}$ multiple times. In comparison, so-called *straightline* extraction works with a single ordinary run of $\mathcal{A}$, without rewinding. Instead, the extractor is then assumed to know some trapdoor information, or it is given enhanced control over some part of the setting. For instance, in the above construction of an extractable commitment, the extractor is given "read access" to $\mathcal{A}$'s random-oracle queries.

Another binary criterion is whether the extraction takes place *on-the-fly*, i.e., during the run of the protocol, or *after-the-fact*, i.e., at the end of the execution. For instance, in the context of proving CCA security for an encryption scheme, to simulate decryption queries without knowing the secret key, it is necessary to extract the plaintext for a queried ciphertext on-the-fly; otherwise, the attacker may abort and not produce the output for which the reduction is waiting.

The extractability of our running example of an extractable commitment in the ROM is *both*, straightline and on-the-fly; we refer to this combination as *online* extraction. This is what we are