

2018–
2022

arXiv:2107.03410v1 [quant-ph] 7 Jul 2021

Quantum algorithms for multivariate Monte Carlo estimation

Arjan Cornelissen[1] and Sofiene Jerbi[2]

[1]QuSoft – University of Amsterdam
[2]Institute for Theoretical Physics, University of Innsbruck

July 2021

Abstract

We consider the problem of estimating the expected outcomes of Monte Carlo processes whose outputs are described by multidimensional random variables. We tightly characterize the quantum query complexity of this problem for various choices of oracle access to the Monte Carlo process and various normalization conditions on its output variables and the estimates to be made. More specifically, we design quantum algorithms that provide a quadratic speed-up in query complexity with respect to the precision of the resulting estimates and show that these complexities are asymptotically optimal in each of our scenarios. Interestingly, we find that these speed-ups with respect to the precision come at the cost of a (up to exponential) slowdown with respect to the dimension of the Monte Carlo random variables, for most of the parameter settings we consider. We also describe several applications of our algorithms, notably in the field of machine learning, and discuss the implications of our results.

1 Introduction

Monte Carlo methods are used extensively in various fields of science and engineering, such as statistical physics [BCH+12], finance [Gla13], or machine learning [ADFDJ03]. At the core of these methods is a Monte Carlo process, e.g., a randomized algorithm $\mathcal{A}$, whose expected outcome is to be estimated via repeated random executions.

An interesting question in complexity theory is whether, given the ability to simulate a Monte Carlo process $\mathcal{A}$ on a quantum computer, quantum algorithms can reduce the sample complexity of such Monte Carlo estimations, i.e., the number of executions of $\mathcal{A}$ needed to obtain a close estimate. From this perspective, the univariate case where $\mathcal{A}$ produces a single real-valued outcome $v(\mathcal{A}) \in [-B, B] \subset \mathbb{R}$ is well studied. An ε-close estimate (in additive error) of $\mathbb{E}_{\mathcal{A}}[v(\mathcal{A})]$ can indeed be estimated with high probability using $\widetilde{\mathcal{O}}\left(\frac{B^2}{\varepsilon^2}\right)$ executions of $\mathcal{A}$ classically, while $\widetilde{\mathcal{O}}\left(\frac{B}{\varepsilon}\right)$ simulations of $\mathcal{A}$ are sufficient quantumly [Mon15].[1] Moreover, these sample complexities are known to be asymptotically optimal, up to polylogarithmic factors, since estimating mean values reduces to this problem [CEG95, NW99].

The multivariate case $v(\mathcal{A}) \in \mathbb{R}^d$ however, appearing notably in machine learning applications (see Section 6), remains largely unaddressed by quantum algorithms. Classically, the multivariate problem does not introduce any special considerations: the d additive approximations of $\mathbb{E}_{\mathcal{A}}[v_1(\mathcal{A})], \ldots, \mathbb{E}_{\mathcal{A}}[v_d(\mathcal{A})]$ can all be computed simultaneously (i.e., using the same executions of $\mathcal{A}$) with only a logarithmic overhead $\log(d)$ in sample complexity (due to the Hoeffding bound, see Appendix A). In the quantum scenario however, this *simultaneous* evaluation of several expectation values is more complicated. The quantum algorithms for the univariate case rely on quantum amplitude estimation [BHMT02], which involves as a critical step an encoding of the expectation values in the relative phase of a quantum register. Given that the phase is bounded and periodic, this imposes a normalization constraint on the encoding of multiple expectation

[1]The $\widetilde{\mathcal{O}}$ notation hides polylogarithmic factors in B and $1/\varepsilon$.