

# Average-Case Verification of the Quantum Fourier Transform Enables Worst-Case Phase Estimation

Noah Linden[1] and Ronald de Wolf[2]

[1]School of Mathematics, University of Bristol. n.linden@bristol.ac.uk
[2]QuSoft, CWI and University of Amsterdam, the Netherlands. rdewolf@cwi.nl

arXiv:2109.10215v3 [quant-ph] 2 Dec 2022

The quantum Fourier transform (QFT) is a key primitive for quantum computing that is typically used as a subroutine within a larger computation, for instance for phase estimation. As such, we may have little control over the state that is input to the QFT. Thus, in implementing a good QFT, we may imagine that it needs to perform well on arbitrary input states. *Verifying* this worst-case correct behaviour of a QFT-implementation would be exponentially hard (in the number of qubits) in general, raising the concern that this verification would be impossible in practice on any useful-sized system. In this paper we show that, in fact, we only need to have good *average-case* performance of the QFT to achieve good *worst-case* performance for key tasks—phase estimation, period finding and amplitude estimation. Further we give a very efficient procedure to verify this required average-case behaviour of the QFT.

## 1 Introduction

### 1.1 Verification of quantum circuits

Massive efforts are currently being expended around the world on building large quantum computers, in academia and industry. Because of the fragility of quantum hardware and the quantum states it produces, it is crucial to be able to *test* that the hardware works as advertised. Such a test may involve some quantum hardware itself, but should be more "lightweight" than the procedure that is being tested, in order to avoid circularity.

There is an important issue with testing that is sometimes overlooked in high-level discussions of the topic: if the circuit of interest is a subroutine in a larger computation, then we may have little control of the state that is input to it; thus we would like to verify that it works on the *worse-case* input state. However, we can typically only test its behavior for an *average-case* input state, because there are far too many possible input states to test them all. In general, testing worst-case correctness of a given $n$-qubit circuit would take resources that scale exponentially in $n$. This means that efficient verification of worst-case correctness typically requires additional assumptions, ranging from restrictions on the class of circuits one is verifying (for instance Clifford circuits [FL11, dSLCP11, LW21]) to cryptographic assumptions (as in Mahadev's approach [Mah18], which also assumes the computation starts with a fixed initial state). For further discussion and pointers to related work on verification of quantum hardware, we refer to our recent paper [LW21] and to the general survey [EHW+20].[1]

[1]The issue of average-case vs worst-case behavior has also recently received attention in the area of Hamiltonian simulation [ZZS+21, CB21].