

2019–
2023

Improved quantum lower and upper bounds for matrix scaling

Sander Gribling* Harold Nieuwboer†

arXiv:2109.15282v1 [quant-ph] 30 Sep 2021

Abstract

Matrix scaling is a simple to state, yet widely applicable linear-algebraic problem: the goal is to scale the rows and columns of a given non-negative matrix such that the rescaled matrix has prescribed row and column sums. Motivated by recent results on first-order quantum algorithms for matrix scaling, we investigate the possibilities for quantum speedups for classical second-order algorithms, which comprise the state-of-the-art in the classical setting.

We first show that there can be essentially no quantum speedup in terms of the input size in the high-precision regime: any quantum algorithm that solves the matrix scaling problem for $n \times n$ matrices with at most m non-zero entries and with ℓ_2-error $\varepsilon = \widetilde{\Theta}(1/m)$ must make $\widetilde{\Omega}(m)$ queries to the matrix, even when the success probability is exponentially small in n. Additionally, we show that for $\varepsilon \in [1/n, 1/2]$, any quantum algorithm capable of producing $\frac{\varepsilon}{100}$-ℓ_1-approximations of the row-sum vector of a (dense) normalized matrix uses $\Omega(n/\varepsilon)$ queries, and that there exists a constant $\varepsilon_0 > 0$ for which this problem takes $\Omega(n^{1.5})$ queries.

To complement these results we give improved quantum algorithms in the low-precision regime: with quantum graph sparsification and amplitude estimation, a box-constrained Newton method can be sped up in the large-ε regime, and outperforms previous quantum algorithms. For entrywise-positive matrices, we find an ε-ℓ_1-scaling in time $\widetilde{O}(n^{1.5}/\varepsilon^2)$, whereas the best previously known bounds were $\widetilde{O}(n^2\mathrm{polylog}(1/\varepsilon))$ (classical) and $\widetilde{O}(n^{1.5}/\varepsilon^3)$ (quantum).

1 Introduction

The matrix scaling problem asks to scale each row and column of a given matrix $\boldsymbol{A} \in [0,1]^{n\times n}$ by a positive number in such a way that the resulting matrix has marginals (i.e., row- and column-sums) that are close to some prescribed marginals. For example, one could ask to scale the matrix in such a way that it becomes doubly stochastic.

Matrix scaling has applications in a wide variety of areas including numerical linear algebra [ABB+99], optimal transport in machine learning [Cut13], statistics [Kru37, DS40, Bro59, BFH75], and also in more theoretical settings, e.g. for approximating the permanent [LSW00]. For a survey, we refer the reader to [Ide16]. Furthermore, the matrix scaling problem is a special (commutative) instance of a more general (non-commutative) class of problems, which includes operator and tensor scaling; these problems have many more applications and are a topic of much recent interest [GGOW19, BFG+19].

Formally, the matrix scaling problem is defined for the ℓ_p-norm as follows. Given a matrix $\boldsymbol{A} \in [0,1]^{n\times n}$ with at most m non-zero entries, entrywise-positive target marginals $\mathbf{r}, \mathbf{c} \in \mathbb{R}^n$ with

*IRIF, Université de Paris, CNRS, Paris, France. Partially supported by SIRTEQ-grant QuIPP. gribling@irif.fr

†Korteweg–de Vries Institute for Mathematics and QuSoft, University of Amsterdam. Supported by NWO grant OCENW.KLEIN.267. h.a.nieuwboer@uva.nl