

2018–2022

Near-Optimal Quantum Algorithms for Multivariate Mean Estimation

Arjan Cornelissen* Yassine Hamoudi† Sofiene Jerbi‡

July 20, 2022

arXiv:2111.09787v2 [quant-ph] 19 Jul 2022

Abstract

We propose the first near-optimal quantum algorithm for estimating in Euclidean norm the mean of a vector-valued random variable with finite mean and covariance. Our result aims at extending the theory of multivariate sub-Gaussian estimators [LM19a] to the quantum setting. Unlike classically, where any univariate estimator can be turned into a multivariate estimator with at most a logarithmic overhead in the dimension, no similar result can be proved in the quantum setting. Indeed, Heinrich [Hei04] ruled out the existence of a quantum advantage for the mean estimation problem when the sample complexity is smaller than the dimension. Our main result is to show that, outside this low-precision regime, there does exist a quantum estimator that outperforms any classical estimator. More precisely, we prove that the approximation error can be decreased by a factor of about the square root of the ratio between the dimension and the sample complexity. Our approach is substantially more involved than in the univariate setting, where most quantum estimators rely only on phase estimation. We exploit a variety of additional algorithmic techniques such as linear amplitude amplification, the Bernstein-Vazirani algorithm, and quantum singular value transformation. Our analysis is also deeply rooted in proving concentration inequalities for multivariate truncated statistics.

We develop our quantum estimators in two different input models that showed up in the literature before. The first one provides coherent access to the binary representation of the random variable and it encompasses the classical setting. In the second model, the random variable is directly encoded into the phases of quantum registers. This model arises naturally in many quantum algorithms but it is often incomparable to having classical samples. We adapt our techniques to these two settings and we show that the second model is strictly weaker than the other one for solving the mean estimation problem. Finally, we describe several applications of our algorithms, notably in measuring the expectation values of commuting observables and in the field of machine learning.

1 Introduction

Monte Carlo methods are used extensively in various fields of science and engineering, such as statistical physics [BH10], finance [Gla03], or machine learning [AFDJ03]. At the core of these methods is a Monte Carlo process, e.g., a randomized algorithm, whose *expected outcome* is to be estimated via repeated random executions. Quantum computers can speed-up this approach at two different levels [Mon15]. First, novel algorithmic techniques such as Hamiltonian simulation [Fey82] or quantum walks [Sze04] provide faster Monte Carlo simulation processes. Secondly, quantum metrology algorithms (such as phase estimation [Kit95]) give better error rates for computing statistics on these processes. The present paper focuses on this second point through the lens of the *mean estimation problem*. In this problem, the objective is to compute

*QuSoft, University of Amsterdam. arjan.cornelissen@cwi.nl
†Simons Institute for the Theory of Computing, University of California, Berkeley. hamoudi@berkeley.edu
‡Institute for Theoretical Physics, University of Innsbruck. sofiene.jerbi@uibk.ac.at

