

Quantum Subroutine Composition

Stacey Jeffery*

CWI & QuSoft, the Netherlands

February 14, 2025

arXiv:2209.14146v3 [quant-ph] 13 Feb 2025

Abstract

An important tool in algorithm design is the ability to build algorithms from other algorithms that run as subroutines. In the case of quantum algorithms, a subroutine may be called on a superposition of different inputs, which complicates things. For example, a classical algorithm that calls a subroutine Q times, where the average probability of querying the subroutine on input i is $\bar{p}_i$, and the cost of the subroutine on input i is a random variable T_i, incurs expected cost $\mathsf{Q}\sum_i \bar{p}_i\mathbb{E}[T_i]$ from all subroutine queries. While this statement is obvious for classical algorithms, for quantum algorithms, it is much less so, since naively, if we run a quantum subroutine on a superposition of inputs, we need to wait for all branches of the superposition to terminate before we can apply the next operation. We nonetheless show an analogous quantum statement (*): If $\bar{q}_i$ is the average *query weight on* i over all queries, the cost from all quantum subroutine queries is $\mathsf{Q}\sum_i \bar{q}_i\mathbb{E}[T_i]$. Here the query weight on i for a particular query is the probability of measuring i in the input register if we were to measure right before the query.

We prove this result using the technique of multidimensional quantum walks, recently introduced in [Jeffery, Zur 2022]. We present a more general version of their quantum walk edge composition result, which yields *variable-time quantum walks*, generalizing variable-time quantum search, by, for example, replacing the update cost with $\sqrt{\sum_{u,v}\pi_u P_{u,v}\mathbb{E}[T_{u,v}^2]}$, where $T_{u,v}$ is the cost to move from vertex u to vertex v. The same technique that allows us to compose quantum subroutines in quantum walks can also be used to compose in any quantum algorithm, which is how we prove (*).

1 Introduction

Classical algorithms make extensive use of subroutines, which allow an algorithm designer to build up an algorithm in modular blocks. Analysis of such algorithms can also be done in a simple, modular way. Suppose an outer classical algorithm makes use of some subroutine at Q steps of the algorithm, and suppose the subroutine's stopping time on input i is some random variable T_i. Suppose $\bar{p}_i$ is the average probability, over all Q queries to the subroutine, that the subroutine is called on input i – so $\mathsf{Q}\bar{p}_i$ is the expected number of times the subroutine is called on input i, and $\sum_i \bar{p}_i = 1$. Let L denote the total number of *other* operations used by the outer algorithm. Then the expected running time of the classical algorithm is:

$$\mathsf{L} + \mathsf{Q}\sum_i \bar{p}_i\mathbb{E}[T_i]. \tag{1}$$

We can get a bounded-error algorithm with this asymptotic running time by simply stopping after 10 times the expected number of steps, if the algorithm has not already halted. At that point we can output an arbitrary (likely wrong) answer if the algorithm hasn't already halted, but this happens with probability at most some small constant.

While subroutines are also quite useful in quantum algorithms, it is not at all obvious how to do the above classical analysis in a quantum setting. Suppose an outer quantum algorithm makes use of

*This work is supported by ERC STG grant 101040624-ASC-Q, NWO Klein project number OCENW.Klein.061, and ARO contract no W911NF2010327. SJ is a CIFAR Fellow in the Quantum Information Science Program.