


Thrifty shadow estimation: re-using quantum circuits and bounding tails

Jonas Helsen[1] and Michael Walter[2]
[1]*QuSoft and CWI, The Netherlands*
[2]*Faculty of Computer Science, Ruhr University Bochum, Germany*

Shadow estimation is a recent protocol that allows estimating exponentially many expectation values of a quantum state from "classical shadows", obtained by applying random quantum circuits and computational basis measurements. In this paper we study the statistical efficiency of this approach in light of near-term quantum computing. We propose a more practical variant of the protocol, *thrifty shadow estimation*, in which quantum circuits are reused many times instead of having to be freshly generated for each measurement. We show that reuse is maximally effective when sampling Haar random unitaries, and maximally ineffective when sampling from the Clifford group, i.e., one should not reuse circuits when performing shadow estimation with the Clifford group. We provide an efficiently simulable family of quantum circuits that interpolates between these extremes, which we believe should be used instead of the Clifford group. Finally, we consider tail bounds for shadow estimation and discuss when median-of-means estimation can be replaced with standard mean estimation.

arXiv:2212.06240v2 [quant-ph] 24 Dec 2023

A key aspect of the development of larger-scale quantum computers is the availability of protocols that can diagnose errors and noise in quantum computations. Over the years many such protocols have been proposed, optimizing either for informational completeness (various forms of tomography) or sampling efficiency (e.g., randomized benchmarking [1, 2] or direct fidelity estimation [2]), but not achieving both at the same time. See [3] for a general overview. Recently Huang, Küng and Preskill (HKP) went beyond this apparent dichotomy by proposing *shadow estimation* [4], a randomized protocol which extracts *exponentially* many expectation values $\mathrm{Tr}(O\rho)$ from *polynomially* many copies of the state ρ, with the only caveat being a restriction on the set of allowed observables O [5]. Shadow estimation has generated significant interest and led to several theoretical follow-up works [6, 7] and experimental applications [8, 9] (see also [10] for a comprehensive overview). At its face value, the protocol is extremely simple: upon receiving a state ρ one generates a random n-qubit circuit U from a circuit set $\mathbb{U}$, applies it to the state ρ, and then measures in the computational basis, obtaining a bit string x. The tuple (U, x) then forms a so-called *classical shadow* of the state ρ, from which expectation values $\mathrm{Tr}(O\rho)$ can be reconstructed by classical post-processing (see Fig. 1 (a) for details). The performance of the protocol depends on the circuit set $\mathbb{U}$, as well as the observables one considers. An important case is when the circuit set is the multi-qubit Clifford group $\mathbb{C}_n$, which is a 3-design. In this case, shadow tomography is efficient for observables O for which $\mathrm{Tr}(O^2)$ is bounded. The Clifford group furthermore has the advantage that if the observable is, e.g., a projection onto a stabilizer state, then the classical post-processing needed is also *computationally efficient* by the Gottesman-Knill theorem [11], which is very useful in practice.

A key component of the HKP proposal is that *every classical shadow requires an independent random circuit*. This is critical to the mathematical argument for its statistical efficiency, but can be undesirable in practice. Especially in near-term quantum computers, it is preferable in many systems to measure a fixed circuit multiple times to generate a large number of classical shadows. This can already be seen in experimental implementations of shadow estimation such as [9], which reports repeating each circuit $> 10^4$ times, and [8], which reports measuring each circuit 10^3 times. This is likely inspired by experience with randomized benchmarking, which similarly samples random circuits (and where the statistics of repeating circuits is well understood [12, 13]). In this work we systematically study the effect of circuit repetition for shadow tomography, which to the best of our knowledge has not been studied before, using tools from representation theory. We also apply those tools to the question of whether median-of-means estimators, another key component of the HKP proposal, are actually necessary for shadow tomography.

THRIFTY SHADOW ESTIMATION

We introduce *thrifty shadow estimation*, our variant of shadow estimation that re-uses quantum circuits and can be significantly more economic in practice. The standard and thrifty protocols protocols are summarized in Fig. 1.

We write N for the total number of measurements, R for the number of times that a random circuit is re-used (including the first time) and K for the number of batches in the median-of-means estimator. We will assume that N is a multiple of KR throughout this section. Thus, the protocol uses N/R random quantum circuits. Note that thrifty shadow estimation reduces to ordinary shadow estimation for $R = 1$. To analyze its statistical performance, note that the thrifty estimator $\hat{o}_R$ is the median-of-means estimator for N/R many i.i.d. copies of the random variable

$$\mathbf{X}_R = \frac{1}{R}\sum_{r=1}^{R} \mathbf{X}^{(r)}, \tag{1}$$

where $\mathbf{X}^{(r)} = \mathrm{Tr}\left(O\mathcal{F}^{-1}(\mathbf{U}\left|\mathbf{x}^{(r)}\right\rangle\!\left\langle\mathbf{x}^{(r)}\right|\mathbf{U}^\dagger)\right)$, with $\mathbf{U}$ drawn uniformly at random from the circuit set $\mathbb{U}$, and $\mathbf{x}^{(1)}, \ldots, \mathbf{x}^{(R)}$ drawn i.i.d. from the conditional distribution

$$p(x|U) = \langle x|\, U\rho U^\dagger \,|x\rangle\,.$$