

2023–
2025

arXiv:2305.13362v1 [quant-ph] 22 May 2023

On quantum backpropagation, information reuse, and cheating measurement collapse

Amira Abbas[1,2,3,4], Robbie King[5], Hsin-Yuan Huang[5,6], William J. Huggins[1], Ramis Movassagh[1], Dar Gilboa[1], and Jarrod R. McClean[1*]

[1] *Google Quantum AI, Venice, California 90291, USA*
[2] *University of KwaZulu-Natal, Durban, South Africa*
[3] *Institute of Physics, University of Amsterdam, Science Park 904, 1098 XH Amsterdam, The Netherlands*
[4] *QuSoft, CWI, Science Park 123, 1098 XG Amsterdam, The Netherlands*
[5] *Department of Computing and Mathematical Sciences, Caltech, Pasadena, CA 91125, USA*
[6] *Institute for Quantum Information and Matter, Caltech, Pasadena, CA 91125, USA*

Abstract

The success of modern deep learning hinges on the ability to train neural networks at scale. Through clever reuse of intermediate information, backpropagation facilitates training through gradient computation at a total cost roughly proportional to running the function, rather than incurring an additional factor proportional to the number of parameters – which can now be in the trillions. Naively, one expects that quantum measurement collapse entirely rules out the reuse of quantum information as in backpropagation. But recent developments in shadow tomography, which assumes access to multiple copies of a quantum state, have challenged that notion. Here, we investigate whether parameterized quantum models can train as efficiently as classical neural networks. We show that achieving backpropagation scaling is impossible without access to multiple copies of a state. With this added ability, we introduce an algorithm with foundations in shadow tomography that matches backpropagation scaling in quantum resources while reducing classical auxiliary computational costs to open problems in shadow tomography. These results highlight the nuance of reusing quantum information for practical purposes and clarify the unique difficulties in training large quantum models, which could alter the course of quantum machine learning.

1 Introduction

Computing gradients through backpropagation is crucial to the success of modern deep neural networks. Rather than a naive manifestation of the chain rule to compute gradients, backpropagation leverages white-box knowledge of a computational graph, as well as intermediate values, to asymptotically improve run times [1–5]. Remarkably, computing the gradient of, say, a neural network function with respect to all its parameters, can be done at a total cost roughly proportional to running the function, instead of incurring an additional factor proportional to the number of parameters. This relative scaling, owed to backpropagation, has facilitated the training of very deep networks, with parameter counts now in order of 10^{10} – accompanied with unparalleled empirical success [1, 6–8]. When considering the number of function calls required to compute gradients, backpropagation in classical circuits remains exponentially more efficient with respect to the number of parameters, than the best known algorithms for determining gradients of parameterized quantum circuits – with or without the aid of a quantum computer [9–13]. Nevertheless, the allure of large-scale models inspires the need for efficient training of parameterized quantum models [14–16], which frequently arise in fields like quantum machine learning and quantum chemistry. But if backpropagation scaling cannot be matched, practically reaching overparameterized regimes may be impossible, even before accounting for additional challenges like barren plateaus [17–19]. Since trainability influences the power and applicability of a model, this could radically shift the current trajectory of preferred quantum models.

*jmcclean@google.com