
2020–
2024


2015–
2019





A Quantum Speed-Up for Approximating the Top Eigenvectors of a Matrix

Yanlin Chen* András Gilyén§ Ronald de Wolf¶

arXiv:2405.14765v2 [quant-ph] 14 Nov 2024

Abstract

Finding a good approximation of the top eigenvector of a given $d \times d$ matrix A is a basic and important computational problem, with many applications. We give two different quantum algorithms that, given query access to the entries of a Hermitian matrix A and assuming a constant eigenvalue gap, output a classical description of a good approximation of the top eigenvector: one algorithm with time complexity $\tilde{\mathcal{O}}(d^{1.75})$ and one with time complexity $d^{1.5+o(1)}$ (the first algorithm has a slightly better dependence on the ℓ_2-error of the approximating vector than the second, and uses different techniques of independent interest). Both of our quantum algorithms provide a polynomial speed-up over the best-possible classical algorithm, which needs $\Omega(d^2)$ queries to entries of A, and hence $\Omega(d^2)$ time. We extend this to a quantum algorithm that outputs a classical description of the subspace spanned by the top-q eigenvectors in time $qd^{1.5+o(1)}$. We also prove a nearly-optimal lower bound of $\tilde{\Omega}(d^{1.5})$ on the quantum query complexity of approximating the top eigenvector.

Our quantum algorithms run a version of the classical power method that is robust to certain benign kinds of errors, where we implement each matrix-vector multiplication with small and well-behaved error on a quantum computer, in different ways for the two algorithms. Our first algorithm estimates the matrix-vector product one entry at a time, using a new "Gaussian phase estimation" procedure. Our second algorithm uses block-encoding techniques to compute the matrix-vector product as a quantum state, from which we obtain a classical description by a new time-efficient unbiased pure-state tomography procedure. This procedure uses an essentially optimal number $\mathcal{O}\big(d\log(d)/\varepsilon^2\big)$ of "conditional sample states"; if we have a state-preparation unitary available rather than just copies of the state, then this ε-dependence can be improved further quadratically. Our procedure comes with improved statistical properties and faster runtime compared to earlier pure-state tomography algorithms. We also develop an almost optimal time-efficient process-tomography algorithm for reflections around bounded-rank subspaces, providing the basis for our top-eigensubspace estimation algorithm, and in turn providing a pure-state tomography algorithm that only requires a reflection about the state rather than a state preparation unitary as input.

*QuSoft, CWI, the Netherlands. yanlin.chen@cwi.nl

§HUN-REN Alfréd Rényi Institute of Mathematics, Budapest, Hungary. Funded by the EU's Horizon 2020 Marie Skłodowska-Curie program QuantOrder-891889 and the QuantERA II project HQCC-101017733 in coordination with the national funding organisation NKFIH. gilyen@renyi.hu

¶QuSoft, CWI and University of Amsterdam, the Netherlands. Partially supported by the Dutch Research Council (NWO) through Gravitation-grant Quantum Software Consortium, 024.003.037. rdewolf@cwi.nl