

Space-Efficient Quantum Error Reduction without log Factors

Aleksandrs Belovs* and Stacey Jeffery†

*Center for Quantum Computing Science, Faculty of Science and Technology, University of Latvia

†QuSoft, CWI & University of Amsterdam, the Netherlands

arXiv:2502.09249v2 [quant-ph] 13 Dec 2025

Abstract

Given an algorithm that outputs the correct answer with bounded error, say 1/3, it is sometimes desirable to reduce this error to some arbitrarily small ε – for example, if one wants to call the algorithm many times as a subroutine. The usual method, for both quantum and randomized algorithms, is a procedure called majority voting, which incurs a multiplicative overhead of $\mathcal{O}(\log \frac{1}{\varepsilon})$ from calling the algorithm this many times.

A recent paper [8] introduced a model of quantum computation called *transducers*, and showed how to reduce the "error" of a transducer arbitrarily with only constant overhead, using a construction analogous to majority voting called *purification*. Even error-free transducers map to bounded-error quantum algorithms, so this does not let you reduce algorithmic error for free, but it does allow bounded-error quantum algorithms to be composed without incurring log factors.

In this paper, we present a new highly simplified construction of a purifier, that can be understood as a weighted walk on a line similar to a random walk interpretation of majority voting. Our purifier has much smaller space and time complexity than the previous one. Indeed, it only uses one additional counter, and only performs two increment and two decrement operations on each iteration. It also has quadratically better dependence on the soundness-completeness gap of the algorithm being purified. We prove that our purifier has optimal query complexity, even down to the constant, which matters when one composes quantum algorithms to super-constant depth.

Purifiers can be seen as a way of turning a "Monte Carlo" quantum algorithm into a "Las Vegas" quantum algorithm – a process for which there is no classical analogue – leading to phenomena like the above-mentioned composition of bounded-error quantum algorithms without log factors. Conceptually, our simplified construction sheds light on this strange quantum phenomenon, and, more practically, could have implications for the time and space complexity of composed quantum algorithms.

1 Introduction

Error reduction is one of the most fundamental algorithmic primitives in both the randomized and quantum settings. It transforms a program that evaluates a function f with bounded error, say 1/3, into a program that evaluates the same function with much smaller error ε. In both the quantum and randomized cases, the conventional procedure is based on majority voting.

If we want to reduce the error of a complete quantum algorithm that ends with a measurement outputting a classical string, then we can treat it as a randomized process and apply majority voting just as we would in a randomized algorithm: execute the quantum algorithm together with the final measurement a number of times, and output the majority.

Somewhat more interesting is error reduction of a *coherent* quantum subroutine. If the subroutine is executed several times, such as when it is employed as an input oracle for another quantum algorithm, then its error must be driven down in order to keep the total error of the