

Predicting Module-Lattice Reduction

Léo Ducas[1,2], Lynn Engelberts[1,3], and Paola de Perthuis[1]

[1] Centrum Wiskunde & Informatica, the Netherlands
[2] Leiden University, the Netherlands
[3] QuSoft, the Netherlands

arXiv:2510.10540v1 [cs.CR] 12 Oct 2025

Abstract. Is module-lattice reduction better than unstructured lattice reduction? This question was highlighted as 'Q8' in the Kyber NIST standardization submission (Avanzi et al., 2021), as potentially affecting the concrete security of Kyber and other module-lattice-based schemes. Foundational works on module-lattice reduction (Lee, Pellet-Mary, Stehlé, and Wallet, ASIACRYPT 2019; Mukherjee and Stephens-Davidowitz, CRYPTO 2020) confirmed the existence of such module variants of LLL and block-reduction algorithms, but focus only on provable worst-case asymptotic behavior.

In this work, we present a concrete average-case analysis of module-lattice reduction. Specifically, we address the question of the expected slope after running module-BKZ, and pinpoint the discriminant Δ_K of the number field at hand as the main quantity driving this slope. We convert this back into a gain or loss on the blocksize β: module-BKZ in a number field K of degree d requires an SVP oracle of dimension $\beta + \ln(|\Delta_K|/d^d)\beta/(d\ln\beta) + o(\beta/\ln\beta)$ to reach the same slope as unstructured BKZ with blocksize β. This asymptotic summary hides further terms that we predict concretely using experimentally verified heuristics. Incidentally, we provide the first open-source implementation of module-BKZ for some cyclotomic fields.

For power-of-two cyclotomic fields, we have $|\Delta_K| = d^d$, and conclude that module-BKZ requires a blocksize larger than its unstructured counterpart by $d-1+o(1)$. On the contrary, for all other cyclotomic fields we have $|\Delta_K| < d^d$, so module-BKZ provides a sublinear $\Theta(\beta/\ln\beta)$ gain on the required blocksize, yielding a subexponential speedup of $\exp(\Theta(\beta/\ln\beta))$.

1 Introduction

Module lattices were introduced in cryptography in 1996 with the NTRU cryptosystem [HPS98], and have since seen increasing interest driven by foundational results on their average-case hardness [SSTX09, LPR10, SS11, LS15]. In particular, they now underlie the security of three NIST post-quantum standards, ML-KEM, ML-DSA, and FN-DSA [NIS22], as well as a plethora of variants.

The central cryptanalytic tool to attack those cryptosystems is block lattice reduction, a term covering a variety of algorithms (including BKZ [Sch87, GN08b], slide [GN08a], and DBKZ [MW16]) that generalize LLL [LLL82] and offer a