

Classical shadows with arbitrary group representations

Maxwell West,[1,2] Frédéric Sauvage,[3] Aniruddha Sen,[4,1] Roy Forestano,[5] David Wierichs,[6] Nathan Killoran,[6] Dmitry Grinko,[7,8,9] M. Cerezo,[10,11] and Martín Larocca[1,11]

[1] *Theoretical Division, Los Alamos National Laboratory, Los Alamos, New Mexico 87545, USA*
[2] *School of Physics, University of Melbourne, Parkville, VIC 3010, Australia*
[3] *Quantinuum, Partnership House, Carlisle Place, London SW1P 1BX, United Kingdom*
[4] *Department of Computer Science, University of Texas at Austin, Austin, TX, 78712, USA*
[5] *Physics Department, Institute for Fundamental Theory, University of Florida, Gainesville, FL, 32611, USA*
[6] *Xanadu, Toronto, ON, M5G 2C8, Canada*
[7] *QuSoft, Amsterdam, The Netherlands*
[8] *Institute for Logic, Language and Computation, University of Amsterdam, Amsterdam, The Netherlands*
[9] *Korteweg-de Vries Institute for Mathematics, University of Amsterdam, Amsterdam, The Netherlands*
[10] *Information Sciences, Los Alamos National Laboratory, Los Alamos, New Mexico 87545, USA*
[11] *Quantum Science Center, Oak Ridge, TN 37931, USA*

arXiv:2604.01429v2 [quant-ph] 14 Oct 2026

Classical shadows (CS) has recently emerged as an important framework to efficiently predict properties of an unknown quantum state. A common strategy in CS protocols is to parametrize the basis in which one measures the state by a random group action; many examples of this have been proposed and studied on a case-by-case basis. In this work, we present a unified theory that allows us to simultaneously understand CS protocols based on sampling from general group representations, extending previous approaches that worked in simplified (multiplicity-free) settings. We identify a class of measurement bases which we call "centralizing bases" that allows us to analytically characterize and invert the measurement channel, minimizing classical post-processing costs. We complement this analysis by deriving general bounds on the sample-complexity necessary to obtain estimates of a given precision. Beyond its unification of previous CS protocols, our method allows us to readily generate new protocols based on other groups, or different representations of previously considered ones. For example, we characterize novel shadow protocols based on sampling from the spin and tensor representations of $\mathsf{SU}(2)$, symmetric and orthogonal groups, and the exceptional Lie group G_2.

I. INTRODUCTION

Learning properties of unknown quantum states is a key component of quantum computation and information, but is significantly hindered by the exponential growth in their complexity as a function of system size. Indeed, this typically renders complete characterization intractable for even the relatively modestly sized systems amenable to experimental control today. To combat this, recent protocols [1–22, 24–30, 32–35] have focused on efficiently learning a small subset of favoured observables, for example few-body [1] or symmetry-respecting [21] operators. In particular, the framework developed in Ref. [1] allows for the prediction of many expectation values from "classical shadows" (CS) of a target state ρ, constructed via measurements made after a random action of some unitary ensemble. Given a desired set of properties, it is preferable to work with a protocol that minimizes the variance of the CS estimator; conversely, given a shadows protocol, a key part of the analysis is the determination of the set of observables that lead to small variances, and thus can be efficiently estimated.

The use of group actions to generate random measurement bases lies at the heart of CS protocols, providing a powerful framework for quantum state estimation. Traditionally, groups such as the local and global Clifford groups [1] – or more recently, matchgate circuits [7, 9] – have been employed to construct these measurement ensembles, each analyzed within its own specialized context. While this case-by-case approach has proven effective, it complicates exploration and experimentation with different groups, as each requires a custom inversion strategy and corresponding variance bounds. An important step in a unifying direction was taken in Ref. [17], where it was shown that if a certain action of the group is multiplicity-free, then one can derive a general expression for the corresponding shadows measurement channel and invert it (a necessary component of the classical postprocessing) analytically, with no significant computational overhead. However, many groups of practical interest do not satisfy this assumption, considerably limiting the applicability of the existing analysis.