


Loop Composition in Quantum Algorithms

Stacey Jeffery[1,2], Manideep Mamindlapally[2], and Alex Baudoin Nguetsa Tankeu[1]

[1]QuSoft & CWI, Amsterdam
[2]University of Amsterdam

arXiv:2605.07518v1 [quant-ph] 8 May 2026

Abstract

The quantum circuit model essentially treats every quantum algorithm as a *straight-line program*. While this view is universal, recent work has shown that it is inconvenient for using different-length quantum subroutines in superposition. Using the quantum walk formalism of quantum algorithms, it is possible to model such *branching* behaviour, and get better composition in this setting.

We apply the above branching composition to Grover's algorithm, which gives a variable-time quantum search algorithm that is *worse* than previous work. The reason it is worse is because branching composition does not take into account another deviation from straight-line programs: looping. We show that by modifying branching composition to also include looping, we can get a complexity that matches previous work. This highlights the importance of properly modeling the program control flow when designing quantum algorithms.

1 Introduction

Any deterministic classical algorithm can be modeled as a *straight-line program*, which is a sequence of operations, applied one after the other, starting from some initial state, and ending in some final state, which should encode the program's output. From this description, it is also clear that a (unitary) quantum circuit, which is the most commonly used model of quantum algorithms, is also a straight-line program, but one where we allow quantum operations, on quantum states.

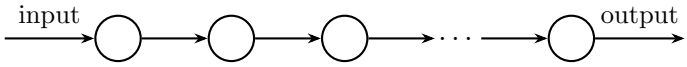


Figure 1: Visualization of a straight-line program.

While the quantum circuit model is universal, one drawback is that it hides basically all structure of the program, aside from its length. One example of where this can give rise to a concrete disadvantage is in subroutine composition. If one of the operations represented by a node in the graph in Figure 1 is something complicated, to be implemented by a subroutine, then this subroutine can also be represented by a straight-line program, which can then be inserted to replace the edge. However, imagine, as is often the case in a quantum algorithm, that this operation is applied to a superposition, and its behaviour is controlled on which branch of the superposition it is running in. One way to think of such a subroutine is as N quantum circuits, specified by sequences of unitaries (i.e. straight-line programs): $\{U^i_{T_i} \dots U^i_1\}_{i=1}^N$, where we want to apply the i-th circuit in the i-th superposition branch. The superposition structure is completely lost in the straight-line program picture. To use this subroutine, we need to express it as a single straight-line program, that calls $U^i_{T_i} \dots U^i_1$ controlled on some particular register containing value $|i\rangle$. Even in the quantum random access model, this represents a straight-line program of length $\max_i T_i$.

To get around this, we need to enhance our straight-line programs to allow them to express superposition branching behaviour, which is quite easy to do visually – see Figure 2 – and somewhat more technical to do mathematically.