


ABSTRACTS: Amsterdam Benchmark Suite for the Time and Resource Analysis of Clifford+T Simulators

Matthew Sutcliffe John van de Wetering
University of Amsterdam, QuSoft
Amsterdam, The Netherlands
m.j.sutcliffe@uva.nl j.m.m.vandewetering@uva.nl

arXiv:2608.24370v1 [quant-ph] 25 Aug 2026

Recent years have seen a rapid growth in literature presenting new methods for simulating non-Clifford quantum circuits with classical hardware. These methods span a range of approaches, including stabiliser decomposition and tensor contraction techniques, varying in efficiency depending on circuit class, depth, non-Clifford gate count, and other metrics. A notable limitation of this literature is the lack of a standardised approach to benchmarking, with each new paper outlining its own specification, simulating its own set of circuits on the authors' own hardware. This paper seeks to address this issue by presenting a standardised and canonical benchmark suite and infrastructure for quantifying the efficiency of non-Clifford classical simulators, with a consistent dataset of circuits and providing consistent (virtual) hardware, thereby enabling a fair comparison of results.

1 Introduction

A core area of research in quantum computing is *classical simulation*: the task of computing quantum circuits with classical hardware. This is vital for verifying quantum algorithms and computers, as well as understanding and quantifying the quantum advantage. Many approaches have been taken in solving this problem, with methods based on state vector simulation [19, 40], tensor contraction [17, 41], stabiliser decomposition [21, 36], and beyond. In the general '*non-Clifford*' case, every known method scales exponentially against one circuit metric or another, such as qubit count, non-Clifford gate count, or treewidth. Many new methods, techniques, and heuristics have been presented in recent years, demonstrating reduced rates of exponential scaling and hence more efficient classical simulation. This in turn enables larger and more complex quantum circuits to be simulable classically and raises the bar quantum computers and algorithms must meet to achieve quantum advantage [18, 31, 30].

With the literature on this topic growing rapidly, a notable problem has arisen, namely the lack of a standardised approach to benchmarking. In general, every new paper presenting a new simulation technique has outlined its own set of benchmark circuits and run experiments on the authors' own hardware, following their own methodology, often benchmarked alongside any one among many existing methods for comparison. While this was sufficient when the literature was small and few competing methods existed, it has since made fair comparisons of simulation techniques difficult.

To address this issue, we present the *Amsterdam Benchmark Suite for the Time and Resource Analysis of Clifford+T Simulators* (*ABSTRACTS*), or simply the *Amsterdam Benchmark Suite*, as a standardised and canonical benchmark dataset and framework for assessing classical simulators of non-Clifford quantum circuits, with an automated execution pipeline as illustrated in Figure 1. This provides consistency in the benchmark setup, methodology, circuits, and hardware spec, as well as a record of results for many existing simulators.

Preprint.

© M. Sutcliffe & J. van de Wetering
This work is licensed under the
Creative Commons Attribution License.