

2022–
2024

Universal entanglement distillation

Salvatore Tirone,[1,2,*] Francesco Anna Mele,[3,†] Vittorio Giovannetti,[4,‡] and Ludovico Lami[5,§]

[1] *QuSoft, Science Park 123, 1098 XG Amsterdam, the Netherlands*
[2] *Korteweg–de Vries Institute for Mathematics, University of Amsterdam, Science Park 105–107, 1098 XG Amsterdam, the Netherlands*
[3] *Institute for Quantum Information and Matter, California Institute of Technology, Pasadena, California 91125, USA*
[4] *NEST, Scuola Normale Superiore and Istituto Nanoscienze, Consiglio Nazionale delle Ricerche, Piazza dei Cavalieri 7, 56126 Pisa, Italy*
[5] *Scuola Normale Superiore, Piazza dei Cavalieri 7, 56126 Pisa, Italy*

We present a universal entanglement distillation protocol that asymptotically achieves the distillable entanglement of any unknown bipartite state in a fixed finite dimension. That is, its ebit yield attains every rate below the optimum that would be available if the state were known, with vanishing trace-distance error. Our approach combines tomography by local operations and classical communication with an adaptive distillation routine, using a local robustness lemma to control finite-sample estimation errors. We also characterize worst-case distillation from a family of states: the optimal uniform rate is the minimum distillable entanglement on the family's closure. These results show that ignorance of the state need not incur an asymptotic rate penalty.

arXiv:2609.40080v1 [quant-ph] 30 Sep 2026

Introduction.—Entanglement is one of the most valuable resources for quantum technologies. Quantum communication, error correction, cryptography, and tests of nonlocality [1–6] rely on entanglement, often supplied in the form of Bell pairs, or ebits. A crucial task in quantum information science is therefore *entanglement distillation* [7, 8], which aims to extract approximate ebits from many identical copies of a bipartite state using only local operations and classical communication (LOCC) [9, 10]. The usual formulation of entanglement distillation assumes that Alice and Bob have a complete classical description of their shared state. In practice, this description may not be available, and finite-sample tomography cannot, in general, determine an arbitrary unknown state exactly. In this work, we therefore study entanglement distillation when the state is partially or entirely unknown. Related state-agnostic questions have been considered in settings ranging from entanglement concentration [11–13], state merging and entanglement distillation [14–16], and quantum communication [17–20] to work extraction [21–27].

We prove the existence of an optimal universal protocol for entanglement distillation from unknown bipartite states. Our approach starts with an LOCC tomography phase, followed by a distillation routine selected on the basis of the resulting estimate. Remarkably, this two-stage strategy asymptotically achieves the full distillable entanglement E_d of every input state with vanishing error. Here, the local dimension is known and the input consists of identical, independently prepared copies; it is the state itself that is unknown. Thus, prior classical knowledge of the state is not a prerequisite for optimal entanglement distillation.

Central to this finding is a technical lemma establishing the local robustness of entanglement distillation. For any distillable state ρ and any rate $0 < R < E_d(\rho)$, we show that there is a sufficiently small open ball around ρ on which a single sequence of LOCC maps achieves rate R, with error tending to zero uniformly over the ball. The radius may depend on both ρ and R. Leveraging this robustness, we also show that the worst-case distillable entanglement over a nonempty family of states $\mathcal{X}$ equals the minimum of E_d on its topological closure $\bar{\mathcal{X}}$.

Conceptually, this framework connects entanglement distillation with quantum learning theory [28–31], and in particular with quantum state tomography [32–35]. While quantum learning theory often focuses on the sample complexity of state characterization, here we use its statistical guarantees to select a resource-distillation protocol. The key point is physically intuitive: the state need not be learned exactly; instead, a sufficiently accurate estimate suffices to identify a robust protocol, and the fraction of copies spent on learning can vanish asymptotically.

Our result thus answers whether a permanent asymptotic rate penalty is an unavoidable toll for ignorance. Under the assumptions above, it is not. The robustness lemma may also be useful in other resource theories, provided that analogous local distillation guarantees can be established.

Preliminaries.—Let us briefly review the setting and the technical tools required for our work. We consider finite-dimensional complex Hilbert spaces, with $\mathcal{H}_A \cong \mathcal{H}_B \cong \mathcal{H}$ and $\mathcal{H} = \mathbb{C}^d$, where $d \geq 2$ is the local dimension and A and B are held by Alice and Bob, respectively. Given a Hilbert space $\mathcal{H}$, we denote by $\mathcal{D}(\mathcal{H})$ the corresponding set of quantum states, namely the positive semidefinite operators on $\mathcal{H}$ with trace one. A shared state ρ_{AB} thus belongs to $\mathcal{D}(\mathcal{H}_A \otimes \mathcal{H}_B)$; we omit subsystem labels when this does not give rise to any ambiguity. All logarithms are to base two.

We measure the distance between states ρ and