

2022–2025





2020–202



Approximate cloning of structured pure states

Yaroslav Herasymenko,[1,*] Junqiao Lin,[2] Philip Verduyn Lunel,[3] and Dmitry Grinko[4]

[1] *Perimeter Institute for Theoretical Physics, Waterloo, Ontario N2L 2Y5, Canada*
[2] *QuSoft, CWI, Amsterdam, The Netherlands*
[3] *Sorbonne Université, CNRS, LIP6, France*
[4] *QuSoft, University of Amsterdam, The Netherlands*

arXiv:2610.06723v1 [quant-ph] 5 Oct 2026

The no-cloning theorem is a cornerstone result in quantum mechanics that forbids copying of general quantum information. More quantitatively, even given N copies of an unknown state, any channel producing an $(N+1)$-copy state incurs a nonzero trace-distance error on some inputs. Turning this around, one can ask about the sample complexity of approximate $N \to N+1$ cloning: for a desired small error ϵ, what N suffices? For arbitrary pure states on a Hilbert space $\mathcal{H}$, the answer is $N \simeq \dim \mathcal{H}/\epsilon$ – astronomically large for most $\mathcal{H}$ of interest. What if the input is promised to lie in a structured family? We examine a range of families fundamental to many-body physics and quantum information: n-mode fermionic Gaussian and Slater states, n-mode bosonic Gaussian states, and qudit phase states. For fermions, we construct optimal cloning channels, reducing the complexity from exponential to polynomial in n. For bosonic Gaussian states, we construct an explicit cloner that certifies polynomial complexity, without any constraint on the energy of the state; this channel, however, is not optimal in general. The fermionic and bosonic results follow from a single representation-theoretic framework, which generalizes previous work by Werner and by Chiribella and Yang; it also yields the cloners' unitary implementation (a Stinespring dilation). Furthermore, we explain why for many families the cloning complexity is linear in the dimension of the family manifold, via a controlled saddle-point evaluation of the frame potential at high N — the same calculation accounts for the analogous scaling of tomography complexity. The above framework does not cover phase states; for these we give an independent construction of a channel with optimal cloning fidelity. For all considered families, our results extend to $N \to N+K$ cloning for $K > 1$. Our work complements recent results on cloning stabilizer states, where cloning via learning yields asymptotically optimal complexity; for our families, we show that such a measure-and-prepare cloner cannot reach small ϵ at any N.

CONTENTS

I. Introduction 3

II. Overview of results 4
A. Obstruction to near-exact cloning by learning 4
B. Gaussian states and resolving orbits 4
C. Stinespring dilation and complementary channels 6
D. Why the manifold dimension controls canonical cloning 7
E. Other examples of resolving orbits 8
F. Optimal cloner for phase states 8
G. Outlook 10

III. Basic notations and definitions 11

IV. Cloning and learning for continuous families 12



* yherasymenko@pitp.ca