



Quickly extracting fidelity decay rates in random circuit benchmarking experiments

Ake Köhne,[1,*] Stavros Efthymiou,[2] Joanna Majsak,[2,3] Jadwiga Wilkens,[4] Bram Evert,[5] Andrew D. Patterson,[5] Jonas Helsen,[1] Markus Heinrich,[3] and Ingo Roth[2]

[1] *CWI & QuSoft, Amsterdam*
[2] *Quantum Research Center, Technology Innovation Institute, Abu Dhabi*
[3] *Institute for Theoretical Physics, University of Cologne, Cologne, Germany*
[4] *Johannes Kepler Universität Linz, Austria*
[5] *Rigetti Computing, 775 Heinz Avenue, Berkeley, California 94710, USA*

(Dated: October 9, 2026)

Randomized benchmarking (RB) with structured circuits is experimentally scalable, but gives biased short-depth fidelity estimates because incomplete ensemble mixing introduces transient modes. Waiting for these transients to decay in deeper circuits can push the signal below the noise floor. This is a well-established obstacle, e.g. in linear cross-entropy benchmarking (LinXEB). Here we overcome this obstacle solely through classical post-processing. We introduce a one-parameter family of estimators that interpolate between the standard LinXEB-type post-processing and a new post-processing technique we call the trace filter function. Under practical assumptions, the trace filter function provably suppresses transient modes, exposing an interpretable fidelity decay at short depth, with theoretical guarantees comparable to standard RB, at the cost of an increased variance. We bound the variance, analyze the evaluation complexity of the new estimators, and develop practical post-processing algorithms. Experiments on a Rigetti superconducting processor validate the method for one-dimensional circuits of up to 12 qubits: where the standard estimator requires depth beyond a 40-layer experimental window, the trace filter function yields consistent estimates below 10 layers. Trace-filtered randomized benchmarking therefore extends reliable fidelity characterization to relevant structured circuit ensembles and short-depth regimes that are inaccessible to conventional estimators.

arXiv:2610.11540v1 [quant-ph] 8 Oct 2026

In the quest to develop large-scale and robust quantum computers, an important role is reserved for protocols aiming to characterize the performance of these devices informing fabrication, calibration, noise mitigation, error correction and algorithmic resource estimation [1–4].

Many of these characterization methods can be grouped under the class of randomized benchmarking (RB) protocols [5–13], which aim at characterizing the average error rate of a gate set. This is done by estimating the decay rate of the implementation fidelity of randomly drawn quantum circuits with increasing depth. Sample-efficiency and robustness to state preparation and measurement (SPAM) errors make RB protocols particularly attractive and widely used in practice.

In practice, standard RB with random Clifford unitaries is limited to very few qubits, typically one or two, because of the use of deep circuits. More scalable benchmarks make use of random quantum circuits [8, 10, 13–16]. Most prominently, this includes the linear cross-entropy benchmarking (LinXEB) protocol [10, 17], originally designed for certifying quantum advantages in sampling tasks, if one considers the decay of the circuit-averaged LinXEB fidelity with circuit depth [5, 16].

The assumptions under which RB protocols yield reliable and interpretable benchmarks are well understood [5, 16, 18–21]. For Markovian, time-stationary, and gate-dependent noise, RB signals are expected to be dominated by a single exponential decay whose decay parameter is related to the fidelity of an effective average noise process. This simplification relies on averaging over sequences that explore the relevant state space sufficiently isotropically, so that only isotropic components of the noise remain visible. For random-circuit RB, or more generally for non-uniform gate distributions, the single-decay regime is reached after a minimal sequence length, at which the circuit distribution has mixed sufficiently toward the corresponding uniform ensemble.

Before the circuit distribution has mixed, the RB signal contains additional transient components associated with the convergence of the sampled circuit ensemble. Neglecting these components biases single-exponential estimates of the decay rate, an issue already recognized by comparisons of the NIST and uniformly sampled Clifford protocols [22], and prominently noted in LinXEB [23–26]. One can avoid this bias by fitting only sufficiently deep, well-mixed sequences; however, this is often experimentally infeasible because noise may degrade the signal before the mixing regime is reached.

In LinXEB, finite-depth mixing effects are typically mitigated either by engineering circuit topologies and local gate ensembles for fast mixing [10, 17, 27], or by introducing ad hoc ratio estimators that normalize the measured signal by the corresponding noiseless signal of the same non-uniform circuit ensemble [23, 28]. The latter approach can remove bias arising from the ideal circuit distribution, but its status as an RB estimator is less clear: it is not evident that the ratio signal remains governed by a single interpretable decay parameter related to an average fidelity, under realistic gate-dependent noise. At the same time, slowly mixing circuit